

Federated Learning:

Collaborative ML without centralized training data

Sara Khalifa and Abdelwahed Khamis

May 03, 2023




Outline

- Recap on Machine Learning /Deep Learning
- Centralised Machine Learning
- On device inference /Distributed Learning
- Federated Learning
- Implementing Federated Learning

Recap on ML/DL

- The goal of machine learning/deep learning is to find a **model**, which produces a desired output given a particular input.

Example task	Given input	Desired output
Image classification		8
Next-word-prediction	<i>Looking forward to your <u>?</u></i>	<i>reply</i>

- Deep Learning shows **great performance** on complex tasks:
 - Computer Vision
 - Natural Language Processing
 - Robotics and Internet of Things

Recap on ML/DL

- It learns and represents complex patterns in the input data.

```
model = ML()  
optimizer = SGD(model)  
for e in epochs:
```

```
    for pos in range(batch_number):
```


```
        input, target = get_batch(dataset, pos)
```

```
        pred = model(input)
```

```
        loss = Loss(pred, target) # compare results
```

```
        loss.backward() # calculate updates to model
```

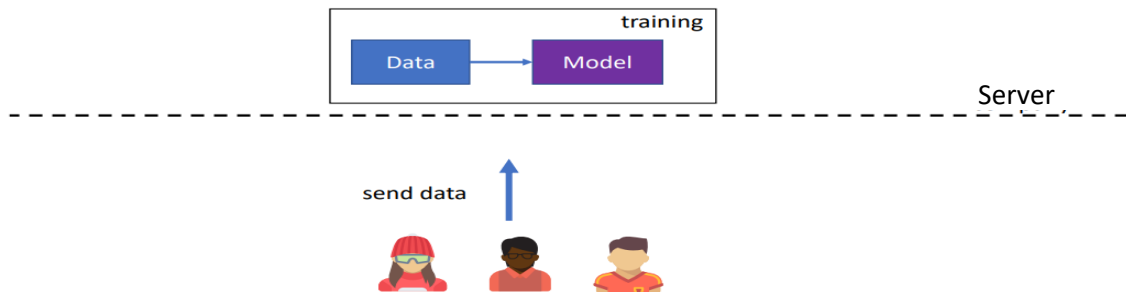
```
        optimizer.step() # apply model
```



Common ML/DL pipeline

Centralized Machine Learning

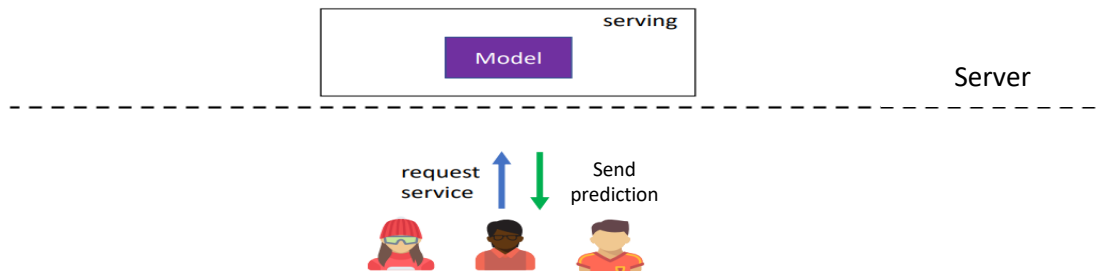
Training



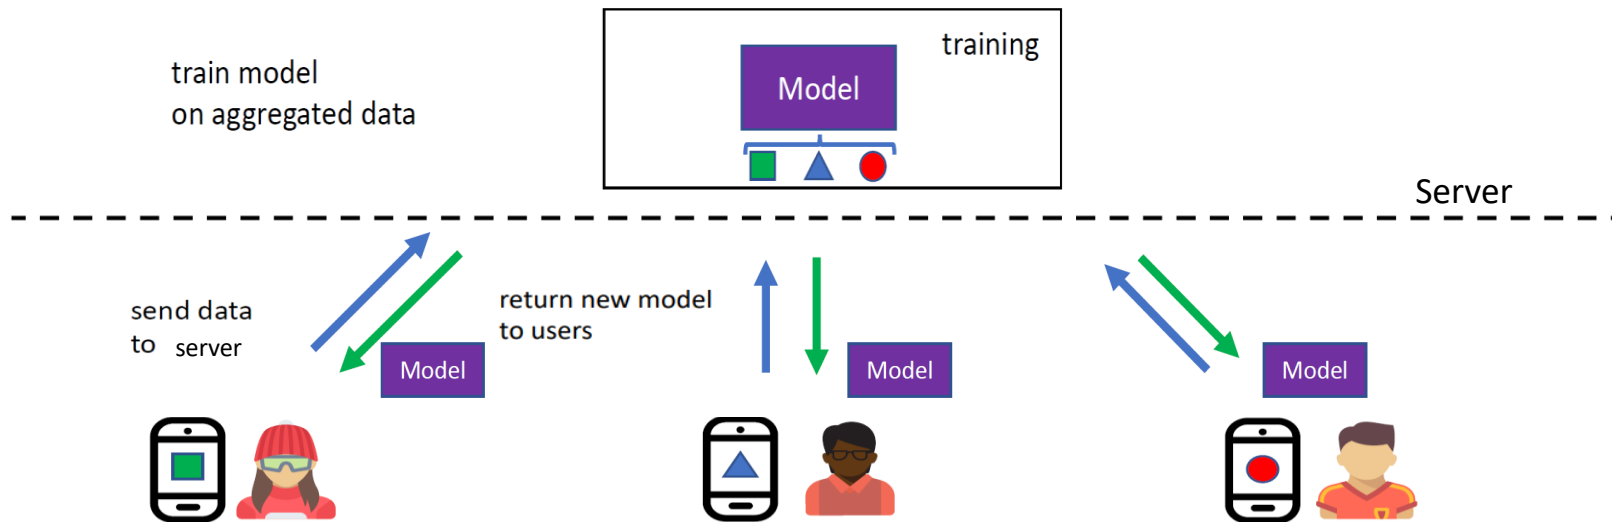
Issues:

- Slow training
- Network Latency
- Limited scalability
- Privacy and security

Testing/Evaluation



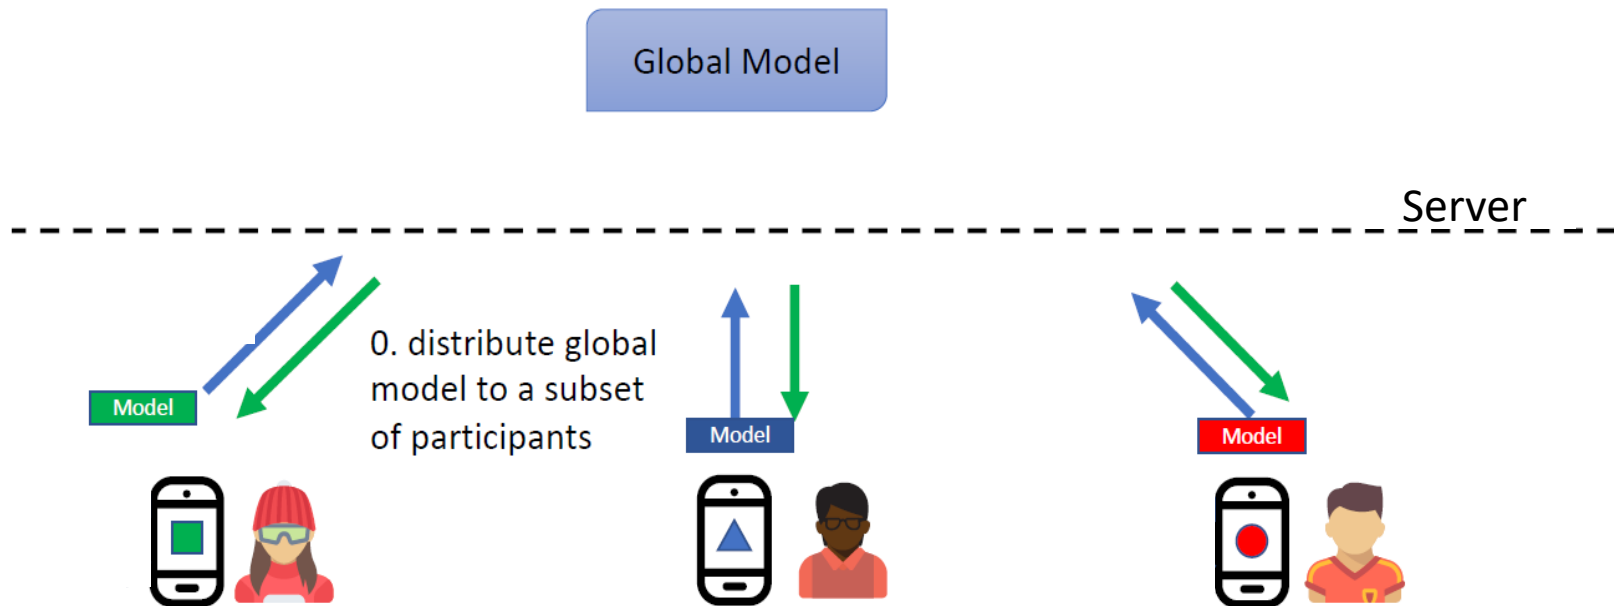
On-device inference / Distributed ML



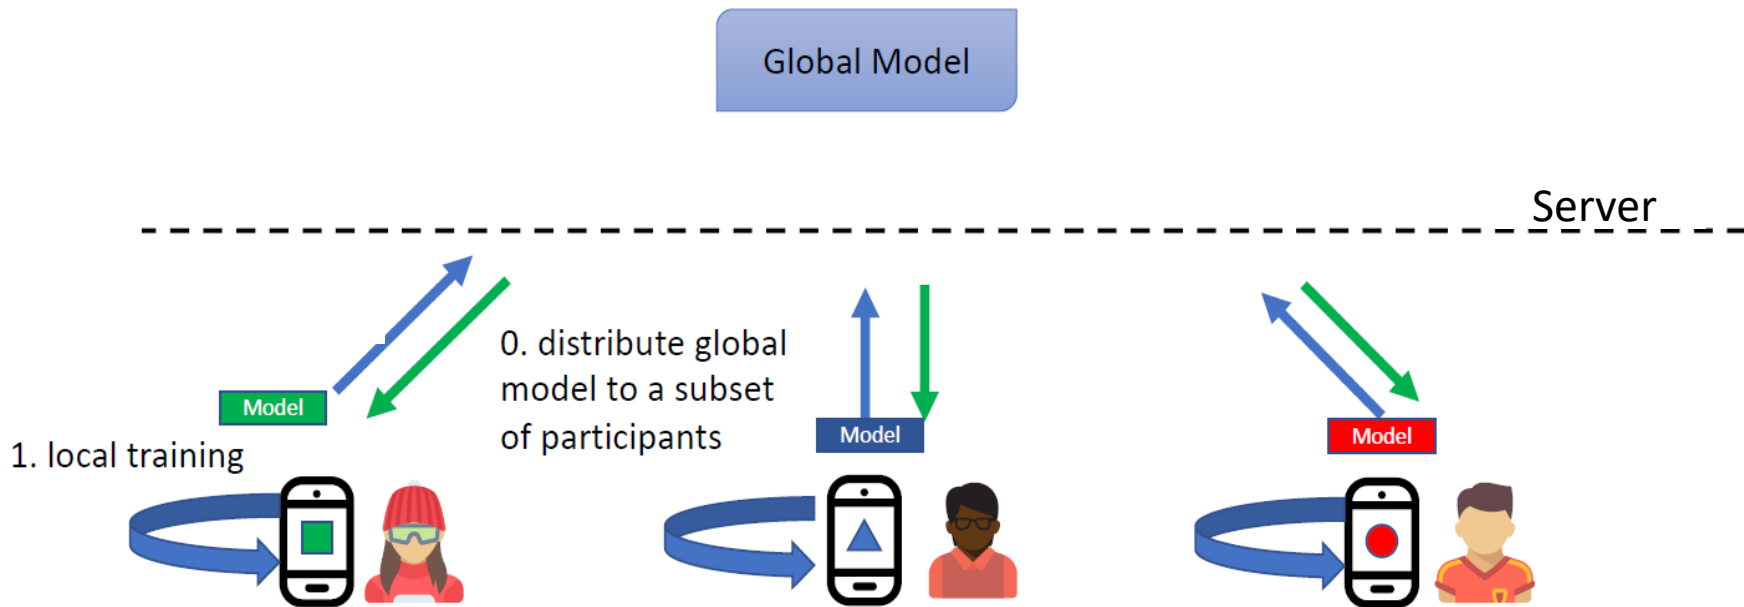
Federated Learning (FL)

- In 2016, Google introduced the concept of FL, enabling collaborative ML **without centralized training data**.
- FL does **not share local data** but ML models, offering applications in diverse domains such as **healthcare, finance, mobile and IoT applications**.
- FL Advantage:
 - Smarter models
 - Reduced latency
 - Less communication cost
 - Privacy preserving

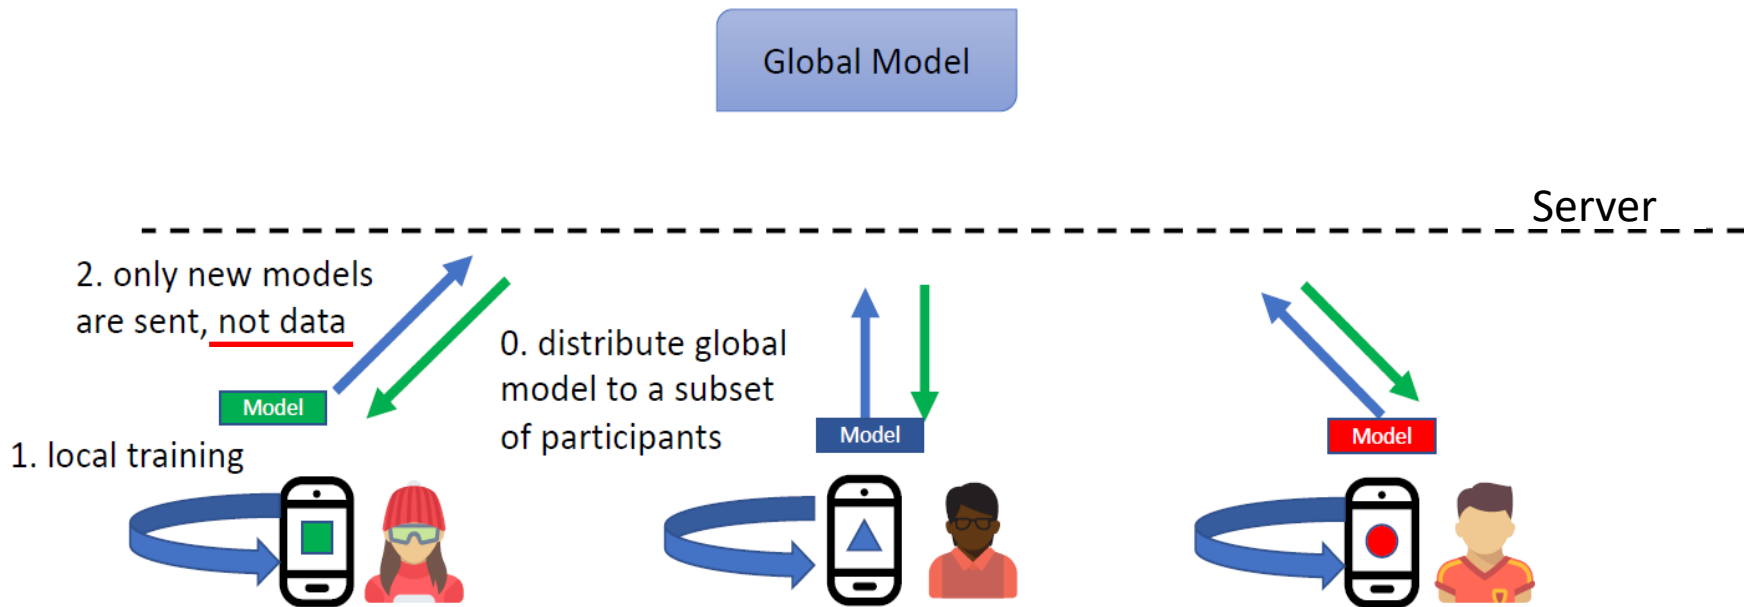
FL Workflow



FL Workflow



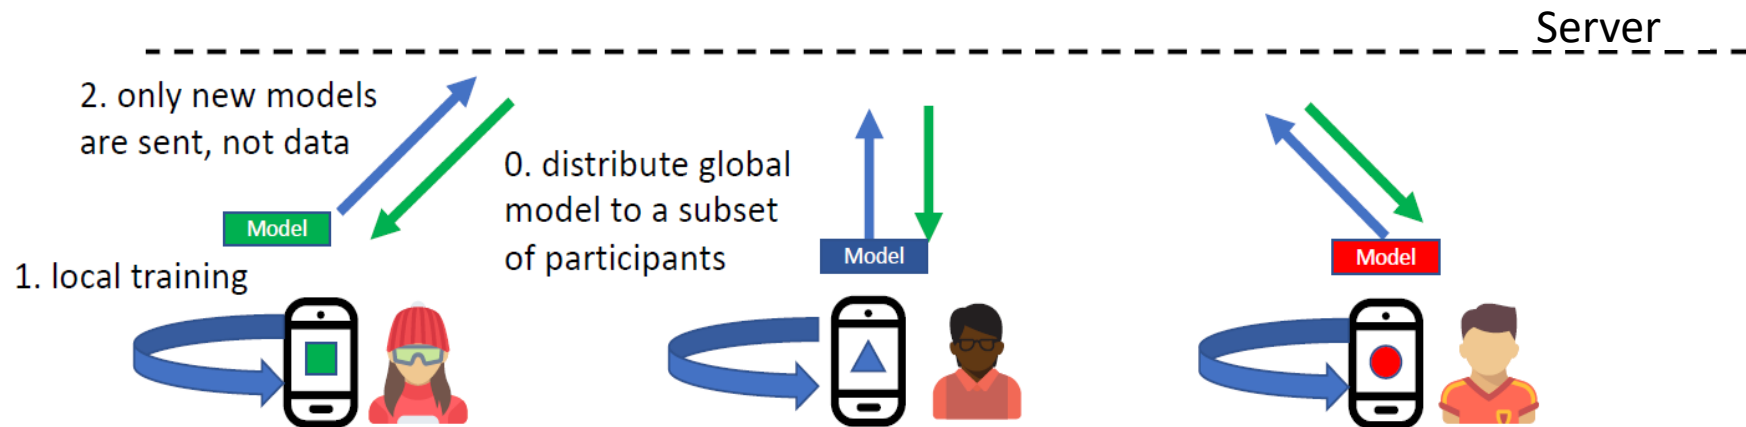
FL Workflow



FL Workflow

3. New models are combined into a global model

Global Model



FL Workflow

3. New models are combined into a global model

Global Model

Server

2. only new models are sent, not data

1. local training

0. distribute global model to a subset of participants



Averaging subset of users allows to train model faster
Algorithm can tolerate drop of participants

User's data distribution and capabilities

- **Non-IID** – every user has different data;
- **Unbalanced** – some users have more data than others;
- **Massively distributed** – millions of smartphones;
- **Limited communication** – mobile devices are frequently offline or on slow or expensive connections.

Model Aggregation

How to aggregate the models after they trained locally?

1. **FedSGD (Federated Stochastic Gradient Descent)**: FedSGD is a simple federated learning algorithm that aggregates the local model updates using **plain averaging**. It does not perform any regularization and assumes that the data distribution across the devices is identically and independently distributed (IID).
2. **FedAvg (Federated Averaging)**: FedAvg is a popular federated learning algorithm that uses a **weighted averaging** scheme to aggregate local model updates based on the number of examples on each device. FedAvg can handle non-IID data distributions and is more robust to noisy updates than FedSGD.
3. **FedAvgM (Federated Averaging with Momentum)**: FedAvgM is an extension of FedAvg that adds **momentum** to the local and global updates to **improve convergence**. FedAvgM can improve the convergence speed and performance compared to FedAvg, especially for non-IID data distributions.
4. **FedProx (Federated Proximal)**: FedProx is a federated learning algorithm that adds a **proximal term** to the local optimization objective to regularize the model and improve **its robustness to noisy updates**.

FedAvg Algorithm

```
1: Server:
2: Initialize global model  $\theta_0$ 
3: for each communication round  $t = 1, 2, \dots T$  do
4:   Select  $m = C \times K$  clients, where  $C \in (0, 1)$ 
5:   for each Client  $k = 1, 2, \dots m$  in parallel do
6:     Download  $\theta_t$  to Client  $k$ 
7:     Do Client  $k$  update and receive  $\theta^k$ 
8:   end for
9:   Update global model  $\theta_t \leftarrow \sum_{k=1}^m \frac{n_k}{n} \theta^k$ 
10: end for
11:
12: Client  $k$  update:
13: Replace local model  $\theta^k \leftarrow \theta_t$ 
14: for local epoch from 1 to  $E$  do
15:   for batch  $b \in (1, B)$  do
16:      $\theta^k \leftarrow \theta^k - \eta \nabla L_k(\theta^k, b)$ 
17:   end for
18: end for
19: Return  $\theta^k$ 
```

T is the total number of communication rounds,

C is the participation ratio assuming that not all local clients participate in each round of model updates

K is the total numbers of clients;

$M = C \times K$ is the number of participating clients

n_k is the local data size, n is the total number of sample pairs

B is the size of mini-batches,

E is the total local training epochs,

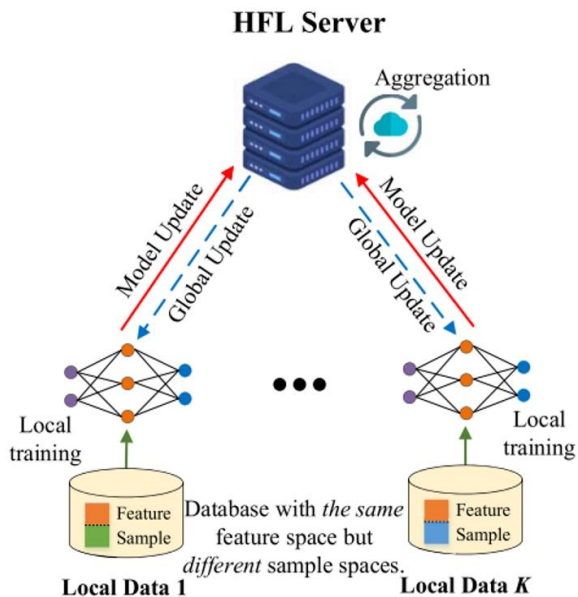
η is the learning rate.

L is the loss function

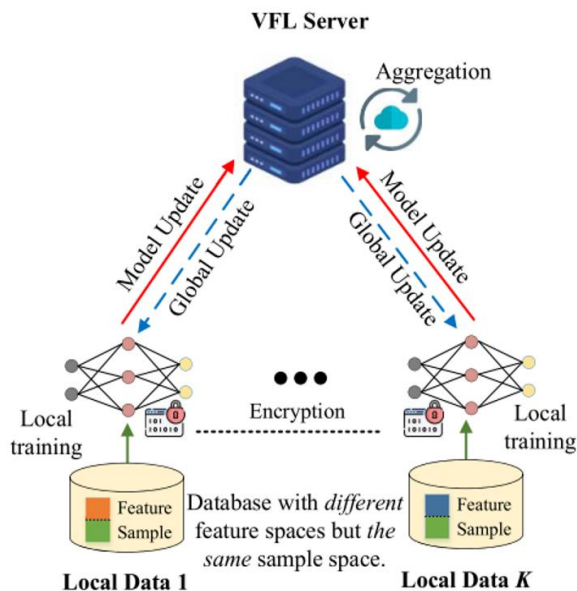
k is the client index

Types of FL

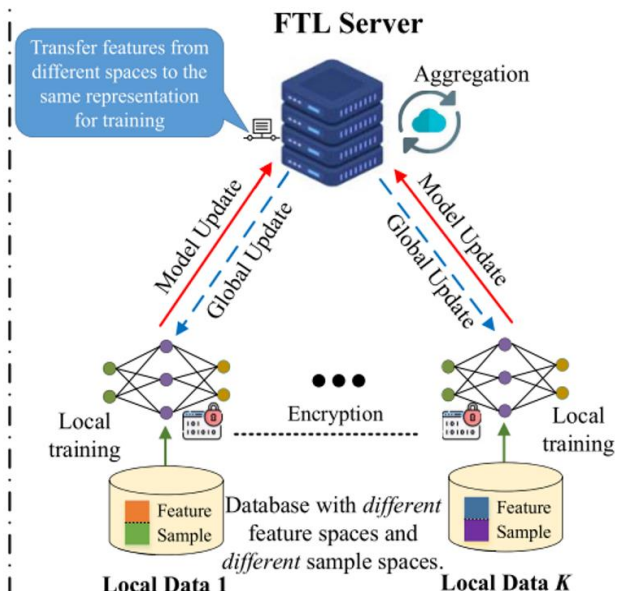
- FL Models based on data partitioning*



(a) Horizontal Federated Learning (HFL)



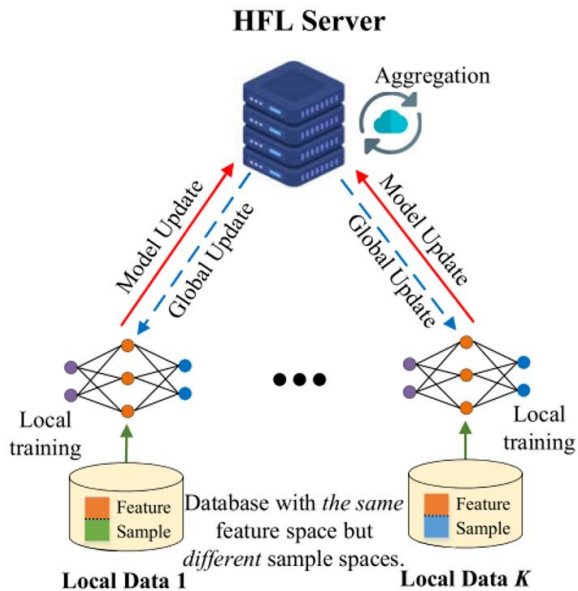
(b) Vertical Federated Learning (VFL)



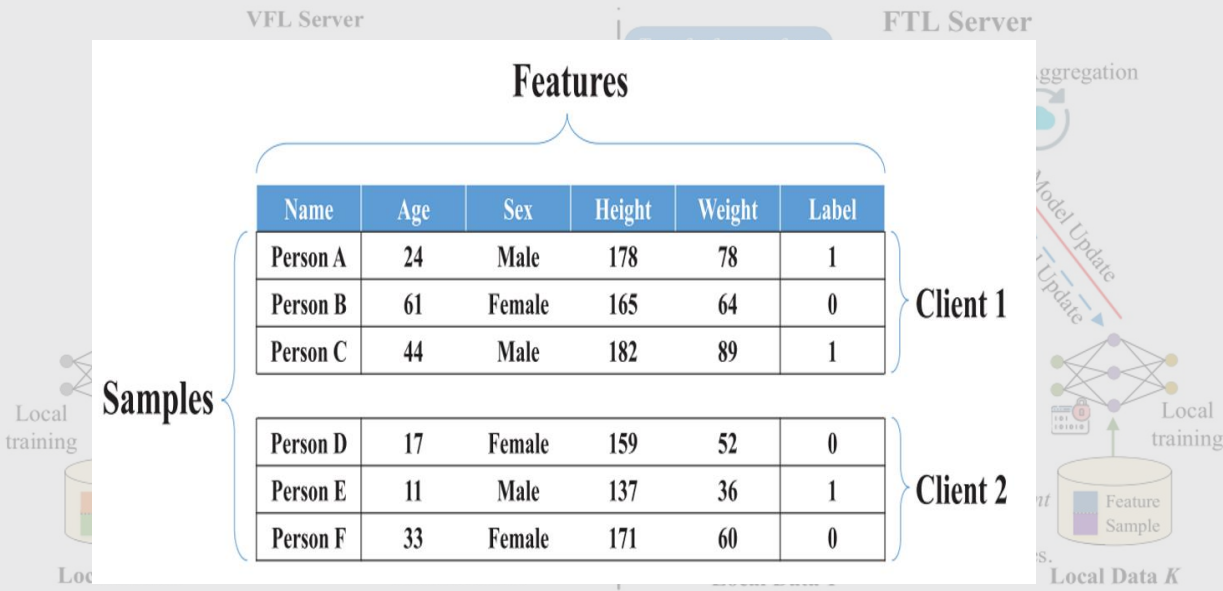
(c) Federated Transfer Learning (FTL)

Types of FL

- FL Models based on data partitioning*



(a) Horizontal Federated Learning (HFL)


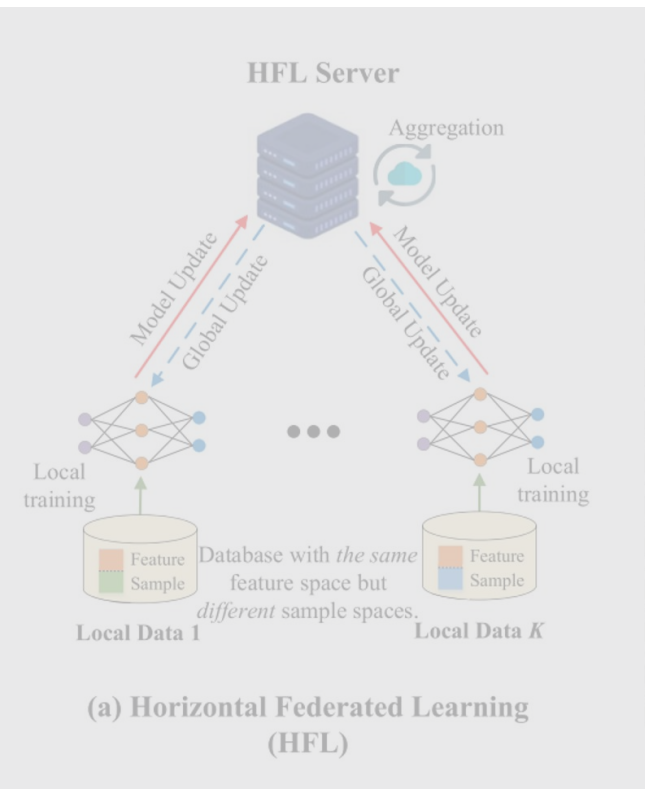


(b) Vertical Federated Learning (VFL)

(c) Federated Transfer Learning (FTL)

Types of FL

- FL Models based on data partitioning**



Features

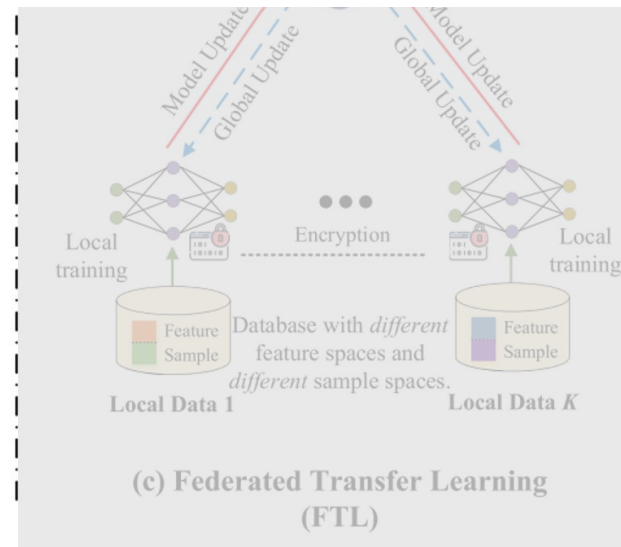
Name	Age	Height	Label
Person A	24	178	1
Person B	61	165	0
Person C	44	182	1
Person D	17	159	0
Person E	11	137	1
Person F	33	171	0

Samples

Name	Sex	Weight
Person A	Male	78
Person B	Female	64
Person C	Male	89
Person D	Female	52
Person E	Male	36
Person F	Female	60

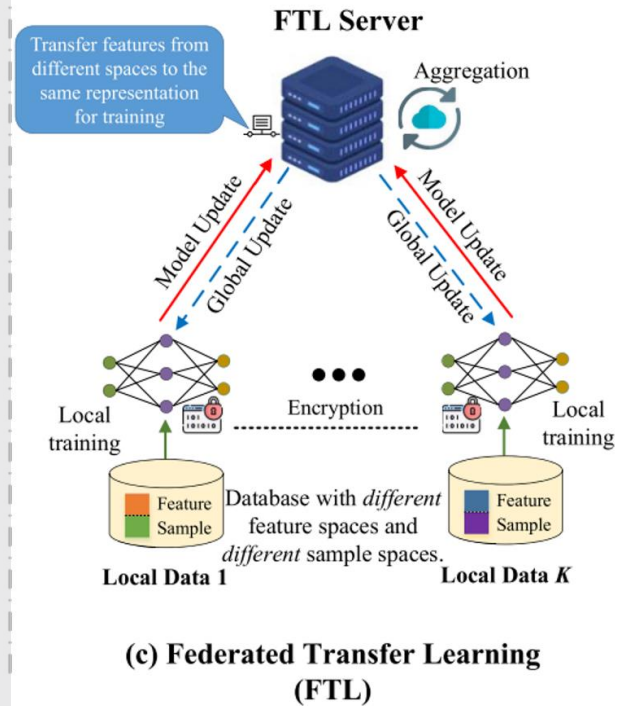
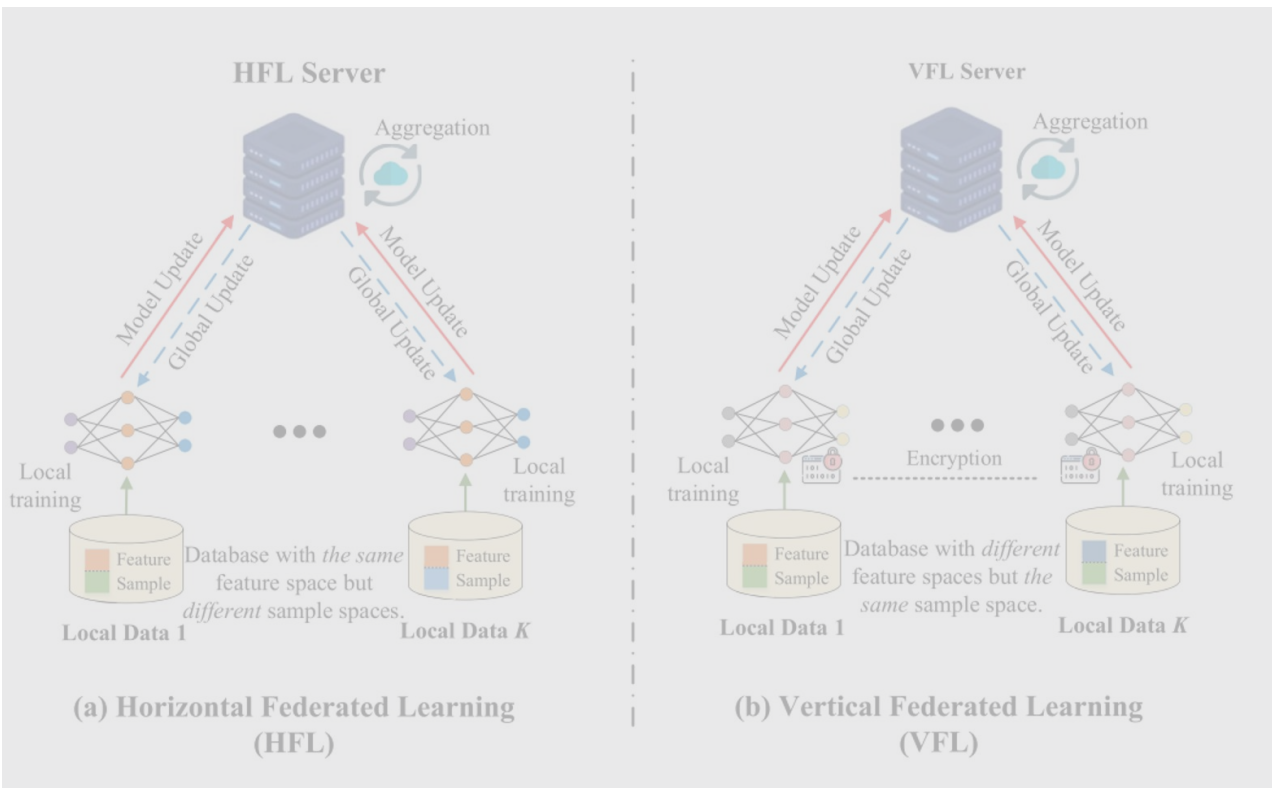
Client 1

Client 2



Types of FL

- FL Models based on data partitioning




Types of FL

- FL Models based on number and nature of clients


Cross-device federated learning

millions of intermittently
available client devices



Cross-silo federated learning

small number of clients
(institutions, data silos),
high availability




Types of FL

- FL Models based on number and nature of clients


Cross-device federated learning

clients cannot be indexed directly (i.e., no use of client identifiers)



Cross-silo federated learning

each client has an identity or name that allows the system to access it specifically




Types of FL

- FL Models based on number and nature of clients

Cross-device federated learning

Server can only access a (possibly biased) random sample of clients on each round.


Large population => most clients only participate once.



Cross-silo federated learning

Most clients participate in every round.

Clients can run algorithms that maintain local state across rounds.



Challenges in Federated learning



expensive communication

- massive, slow networks



privacy concerns

- user privacy constraints



statistical heterogeneity

- unbalanced, non-IID data



systems heterogeneity

- variable hardware, connectivity, etc

Implementing FL

A reference list of popular federated learning repositories.

Name	Repository	License/Stars	Focus
TF Federated	https://github.com/tensorflow/federated	Apache 2.0 / 1.7k	R&D
FedJAX	https://github.com/google/fedjax	Apache 2.0 / 130	Research
Flower	https://github.com/adap/flower	Apache 2.0 / 529	Usability
FedML	https://github.com/FedML-AI/FedML	Apache 2.0 / 839	Research
PySyft	https://github.com/openmined/pysyft	Apache 2.0 / 7.7k	Privacy / R&D
IBM federated-learning-lib	https://github.com/IBM/federated-learning-lib	Custom / 244	Enterprise

Practical Implementation using Flower and Pytorch frameworks.

https://github.com/abdelwahed/FL_tutorial

Thank you

Dr Abdelwahed Khamis

Dr. Sara Khalifa

Distributed Sensing Systems Research Group
Cyber Physical Systems Research Program
Data61, CSIRO

